

Familiarise Yourself with the Hamiltonian Monte Carlo (HMC) Sampling Algorithm

Code to implement HMC is supplied in the module `hmc_lab4.py` and there is a simple demonstration of its usage in the notebook `demo_hmc.ipynb`.

Included in `hmc_lab4.py` are the functions:

- `sample`: The main function, to return samples from some arbitrary distribution over a state space (e.g. a posterior over combined model parameters θ)
- `one_step`: implements a single cycle of HMC (you needn't call this directly)
- `gradient_check`: A useful helper to check your gradient calculations (can also be called automatically from within `sample`). Highly recommended!

Before undertaking any subsequent tasks, you should examine and run the example notebook `demo_hmc.ipynb`, which offers a simple guide on to how to set up and use the `sample` algorithm.

The sampler is called as:

```
sample(x, energy_func, energy_grad, R, L, epsilon0, burn=0, checkgrad=False, args=())
```

Arguments are:

- | | |
|--------------------------|--|
| <code>x</code> | The initial state (parameter values) of θ . In this exercise, any sensible random initialisation should be fine |
| <code>energy_func</code> | The energy (negative log-probability) of the distribution to be sampled from: $-\log P^*(\theta)$ in our notation |
| <code>energy_grad</code> | Gradient of the above energy function with respect to θ |
| <code>R</code> | Number of samples desired (try a few hundred for testing, then perhaps up to 10,000 for quoting results) |
| <code>L</code> | Number of steps to run the Hamiltonian dynamics within each cycle of the sampler ($L = 100$ is probably safe, though for the Gaussian example below, 25 should suffice) |
| <code>epsilon0</code> | Step-size for the dynamics (see below) |
| <code>burn</code> | Number of initial samples to discard as "burn-in" (e.g. you might set this to 10% of <code>R</code>) |
| <code>checkgrad</code> | A flag specifying whether to check the consistency of the error/gradient functions — it is recommended to set this to <code>True</code> ! |
| <code>args</code> | A list containing any additional (fixed) arguments that need to be passed to both <code>energy_func</code> and <code>energy_grad</code> (later, you will need to pass the data inputs and targets) |

Key to the accuracy of the sampler are the "simulation" parameters `L` and `epsilon0`. The former should be set large enough such that the dynamics have time to travel the full range of the distribution for a given step-size `epsilon0`, which itself should be sufficiently small to ensure the simulation does not "diverge". The best settings are rarely obvious without some experimentation!

Here, the general guidance is:

- Set `L` to 100 (as a minimum) for most regression exercises,
- As an exception, setting `L` to 25 for the simple Gaussian example (next) is sufficient,
- Feel free to experiment (and use a smaller `L` when testing),
- Given choice of `L`, `epsilon0` should be set to be as large as possible while keeping the acceptance rate of the sampler relatively high (e.g. no lower than 80%),
- Note that having an acceptance rate of 100% is not likely to be a good thing. That may indicate that `epsilon0` is too small and the state space is not being properly explored,
- One informal approach may therefore be to increase `epsilon0` until the acceptance rate falls below 80%, then "dial it back down a little".